

```

1 import numpy as np
2
3
4 class Softmax(object):
5
6     def __init__(self, dims=[10, 3073]):
7         self.init_weights(dims=dims)
8
9     def init_weights(self, dims):
10         """
11         Initializes the weight matrix of the Softmax classifier.
12         Note that it has shape (C, D) where C is the number of
13         classes and D is the feature size.
14         """
15         self.W = np.random.normal(size=dims) * 0.0001
16
17     def loss(self, X, y):
18         """
19         Calculates the softmax loss.
20
21         Inputs have dimension D, there are C classes, and we operate on
minibatches
22         of N examples.
23
24         Inputs:
25         - X: A numpy array of shape (N, D) containing a minibatch of data.
26         - y: A numpy array of shape (N,) containing training labels; y[i] = c
means
27         that X[i] has label c, where 0 <= c < C.
28
29         Returns a tuple of:
30         - loss as single float
31         """
32
33         # Initialize the loss to zero.
34         loss = 0.0
35
36         # ===== #
37         # YOUR CODE HERE:
38         # Calculate the normalized softmax loss. Store it as the variable
loss.
39         # (That is, calculate the sum of the losses of all the training
40         # set margins, and then normalize the loss by the number of
41         # training examples.)
42         # ===== #
43
44         a = self.W.dot(X.T).T
45
46         i = 0
47         for row in a:
48             row -= np.max(row) #avoid overflow
49             loss += (np.log(np.sum(np.exp(row))) - row[y[i]])
50             i += 1
51
52         loss /= a.shape[0]
53
54         # ===== #
55         # END YOUR CODE HERE
56         # ===== #

```

```

57     return loss
58
59
60 def loss_and_grad(self, X, y):
61     """
62     Same as self.loss(X, y), except that it also returns the gradient.
63
64     Output: grad -- a matrix of the same dimensions as W containing
65             the gradient of the loss with respect to W.
66     """
67
68     # Initialize the loss and gradient to zero.
69     loss = 0.0
70     grad = np.zeros_like(self.W)
71
72     # ===== #
73     # YOUR CODE HERE:
74     # Calculate the softmax loss and the gradient. Store the gradient
75     # as the variable grad.
76     # ===== #
77
78     a = self.W.dot(X.T).T
79
80     i = 0
81     for row in a:
82         row -= np.max(row) #avoid overflow
83         a_row = np.sum(np.exp(row))
84         loss += (np.log(a_row) - row[y[i]])
85
86         for j in np.arange(self.W.shape[0]):
87             grad[j] += (np.exp(row[j])/a_row) * X[i]
88             grad[y[i]] -= X[i]
89         i += 1
90
91     loss /= a.shape[0]
92     grad /= a.shape[0]
93
94     # ===== #
95     # END YOUR CODE HERE
96     # ===== #
97
98     return loss, grad
99
100 def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
101     """
102     sample a few random elements and only return numerical
103     in these dimensions.
104     """
105
106     for i in np.arange(num_checks):
107         ix = tuple([np.random.randint(m) for m in self.W.shape])
108
109         oldval = self.W[ix]
110         self.W[ix] = oldval + h # increment by h
111         fxph = self.loss(X, y)
112         self.W[ix] = oldval - h # decrement by h
113         fxmh = self.loss(X,y) # evaluate f(x - h)
114         self.W[ix] = oldval # reset
115
116         grad_numerical = (fxph - fxmh) / (2 * h)

```

```

117     grad_analytic = your_grad[ix]
118     rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical)
+ abs(grad_analytic))
119     print('numerical: %f analytic: %f, relative error: %e' %
(grad_numerical, grad_analytic, rel_error))
120
121 def fast_loss_and_grad(self, X, y):
122     """
123     A vectorized implementation of loss_and_grad. It shares the same
124     inputs and outputs as loss_and_grad.
125     """
126     loss = 0.0
127     grad = np.zeros(self.W.shape) # initialize the gradient as zero
128
129     # ===== #
130     # YOUR CODE HERE:
131     # Calculate the softmax loss and gradient WITHOUT any for loops.
132     # ===== #
133
134     a = self.W.dot(X.T).T
135     num_train = a.shape[0]
136
137     a -= np.max(a, axis=1, keepdims=True)
138     a_exp = np.exp(a)
139
140     probs = a_exp / np.sum(a_exp, axis=1, keepdims=True)
141     probs_row = probs[range(num_train), y]
142     probs_log = -np.log(probs_row)
143
144     loss = np.sum(probs_log) / num_train
145
146     probs[range(num_train), y] -= 1
147     grad = (probs.T.dot(X)) / num_train
148
149     # ===== #
150     # END YOUR CODE HERE
151     # ===== #
152
153     return loss, grad
154
155 def train(self, X, y, learning_rate=1e-3, num_iters=100,
156         batch_size=200, verbose=False):
157     """
158     Train this linear classifier using stochastic gradient descent.
159
160     Inputs:
161     - X: A numpy array of shape (N, D) containing training data; there are N
162         training samples each of dimension D.
163     - y: A numpy array of shape (N,) containing training labels; y[i] = c
164         means that X[i] has label 0 ≤ c < C for C classes.
165     - learning_rate: (float) learning rate for optimization.
166     - num_iters: (integer) number of steps to take when optimizing
167     - batch_size: (integer) number of training examples to use at each step.
168     - verbose: (boolean) If true, print progress during optimization.
169
170     Outputs:
171     A list containing the value of the loss function at each training
172     iteration.
173     """
174     num_train, dim = X.shape

```

```

174     num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is
number of classes
175
176     self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the
weights of self.W
177
178     # Run stochastic gradient descent to optimize W
179     loss_history = []
180
181     for it in np.arange(num_iters):
182         X_batch = None
183         y_batch = None
184
185         # ===== #
186         # YOUR CODE HERE:
187         #     Sample batch_size elements from the training data for use in
188         #     gradient descent. After sampling,
189         #     - X_batch should have shape: (dim, batch_size)
190         #     - y_batch should have shape: (batch_size,)
191         #     The indices should be randomly generated to reduce correlations
192         #     in the dataset. Use np.random.choice. It's okay to sample with
193         #     replacement.
194         # ===== #
195
196         indices = np.random.choice(X.shape[0], batch_size)
197         X_batch = X[indices]
198         y_batch = y[indices]
199
200         # ===== #
201         # END YOUR CODE HERE
202         # ===== #
203
204         # evaluate loss and gradient
205         loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
206         loss_history.append(loss)
207
208         # ===== #
209         # YOUR CODE HERE:
210         #     Update the parameters, self.W, with a gradient step
211         # ===== #
212
213         self.W -= learning_rate * grad
214
215         # ===== #
216         # END YOUR CODE HERE
217         # ===== #
218
219         if verbose and it % 100 == 0:
220             print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
221
222     return loss_history
223
224 def predict(self, X):
225     """
226     Inputs:
227     - X: N x D array of training data. Each row is a D-dimensional point.
228
229     Returns:
230     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
231       array of length N, and each element is an integer giving the predicted

```

```
232     class.  
233     """  
234     y_pred = np.zeros(X.shape[1])  
235     # ===== #  
236     # YOUR CODE HERE:  
237     #   Predict the labels given the training data.  
238     # ===== #  
239  
240     a = self.W.dot(X.T).T  
241     y_pred = np.argmax(a, axis=1)  
242  
243     # ===== #  
244     # END YOUR CODE HERE  
245     # ===== #  
246  
247     return y_pred  
248  
249
```