

```

1 import numpy as np
2
3 from .layers import *
4 from .layer_utils import *
5
6
7 class TwoLayerNet(object):
8     """
9     A two-layer fully-connected neural network with ReLU nonlinearity and
10     softmax loss that uses a modular layer design. We assume an input dimension
11     of D, a hidden dimension of H, and perform classification over C classes.
12
13     The architecture should be affine - relu - affine - softmax.
14
15     Note that this class does not implement gradient descent; instead, it
16     will interact with a separate Solver object that is responsible for running
17     optimization.
18
19     The learnable parameters of the model are stored in the dictionary
20     self.params that maps parameter names to numpy arrays.
21     """
22
23     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
24                 dropout=0, weight_scale=1e-3, reg=0.0):
25         """
26         Initialize a new network.
27
28         Inputs:
29         - input_dim: An integer giving the size of the input
30         - hidden_dims: An integer giving the size of the hidden layer
31         - num_classes: An integer giving the number of classes to classify
32         - dropout: Scalar between 0 and 1 giving dropout strength.
33         - weight_scale: Scalar giving the standard deviation for random
34           initialization of the weights.
35         - reg: Scalar giving L2 regularization strength.
36         """
37         self.params = {}
38         self.reg = reg
39
40         # ===== #
41         # YOUR CODE HERE:
42         #   Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
43         #   self.params['W2'], self.params['b1'] and self.params['b2']. The
44         #   biases are initialized to zero and the weights are initialized
45         #   so that each parameter has mean 0 and standard deviation
46         #   weight_scale.
47         #   The dimensions of W1 should be (input_dim, hidden_dim) and the
48         #   dimensions of W2 should be (hidden_dims, num_classes)
49         # ===== #
50         self.params['W1'] = weight_scale * np.random.randn(input_dim,
51 hidden_dims) + 0
52         self.params['W2'] = weight_scale * np.random.randn(hidden_dims,
53 num_classes) + 0
54
55         self.params['b1'] = np.zeros((hidden_dims, 1))
56         self.params['b2'] = np.zeros((num_classes, 1))
57
58         # ===== #

```

```

57     # END YOUR CODE HERE
58     # ===== #
59
60 def loss(self, X, y=None):
61     """
62     Compute loss and gradient for a minibatch of data.
63
64     Inputs:
65     - X: Array of input data of shape (N, d_1, ..., d_k)
66     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
67
68     Returns:
69     If y is None, then run a test-time forward pass of the model and return:
70     - scores: Array of shape (N, C) giving classification scores, where
71       scores[i, c] is the classification score for X[i] and class c.
72
73     If y is not None, then run a training-time forward and backward pass and
74     return a tuple of:
75     - loss: Scalar value giving the loss
76     - grads: Dictionary with the same keys as self.params, mapping parameter
77       names to gradients of the loss with respect to those parameters.
78     """
79     scores = None
80
81     # ===== #
82     # YOUR CODE HERE:
83     #   Implement the forward pass of the two-layer neural network. Store
84     #   the class scores as the variable 'scores'. Be sure to use the layers
85     #   you prior implemented.
86     # ===== #
87
88     out_l1, cache_l1 = affine_forward(X, self.params['W1'],
self.params['b1'])
89     out_relu, cache_relu = relu_forward(out_l1)
90     scores, cache_l2 = affine_forward(out_relu, self.params['W2'],
self.params['b2'])
91
92     # ===== #
93     # END YOUR CODE HERE
94     # ===== #
95
96     # If y is None then we are in test mode so just return scores
97     if y is None:
98         return scores
99
100    loss, grads = 0, {}
101    # ===== #
102    # YOUR CODE HERE:
103    #   Implement the backward pass of the two-layer neural net. Store
104    #   the loss as the variable 'loss' and store the gradients in the
105    #   'grads' dictionary. For the grads dictionary, grads['W1'] holds
106    #   the gradient for W1, grads['b1'] holds the gradient for b1, etc.
107    #   i.e., grads[k] holds the gradient for self.params[k].
108    #
109    #   Add L2 regularization, where there is an added cost  $0.5 * \text{self.reg} * W^2$ 
110    #   for each W. Be sure to include the 0.5 multiplying factor to
111    #   match our implementation.
112    #
113    #   And be sure to use the layers you prior implemented.
114    # ===== #

```

```

115
116     loss, dx2 = softmax_loss(scores, y)
117     reg_loss = 0.5 * self.reg * (np.linalg.norm(self.params['W1'], 'fro')**2
+ np.linalg.norm(self.params['W2'], 'fro')**2)
118     loss += reg_loss
119
120     dh1, dw2, db2 = affine_backward(dx2, cache_l2)
121     da = relu_backward(dh1, cache_relu)
122     dx1, dw1, db1 = affine_backward(da, cache_l1)
123
124     grads['W1'] = dw1 + self.reg * self.params['W1']
125     grads['b1'] = db1.T
126
127     grads['W2'] = dw2 + self.reg * self.params['W2']
128     grads['b2'] = db2.T
129
130     # ===== #
131     # END YOUR CODE HERE
132     # ===== #
133
134     return loss, grads
135
136
137 class FullyConnectedNet(object):
138     """
139     A fully-connected neural network with an arbitrary number of hidden layers,
140     ReLU nonlinearities, and a softmax loss function. This will also implement
141     dropout and batch normalization as options. For a network with L layers,
142     the architecture will be
143
144     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
145
146     where batch normalization and dropout are optional, and the {...} block is
147     repeated L - 1 times.
148
149     Similar to the TwoLayerNet above, learnable parameters are stored in the
150     self.params dictionary and will be learned using the Solver class.
151     """
152
153     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
154                 dropout=0, use_batchnorm=False, reg=0.0,
155                 weight_scale=1e-2, dtype=np.float32, seed=None):
156         """
157         Initialize a new FullyConnectedNet.
158
159         Inputs:
160         - hidden_dims: A list of integers giving the size of each hidden layer.
161         - input_dim: An integer giving the size of the input.
162         - num_classes: An integer giving the number of classes to classify.
163         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0
164         then
165         the network should not use dropout at all.
166         - use_batchnorm: Whether or not the network should use batch
167         normalization.
168         - reg: Scalar giving L2 regularization strength.
169         - weight_scale: Scalar giving the standard deviation for random
170         initialization of the weights.
171         - dtype: A numpy datatype object; all computations will be performed
172         using
173         this datatype. float32 is faster but less accurate, so you should use

```

```

171         float64 for numeric gradient checking.
172         - seed: If not None, then pass this random seed to the dropout layers.
This
173         will make the dropout layers deterministic so we can gradient check the
174         model.
175         """
176         self.use_batchnorm = use_batchnorm
177         self.use_dropout = dropout > 0
178         self.reg = reg
179         self.num_layers = 1 + len(hidden_dims)
180         self.dtype = dtype
181         self.params = {}
182
183         # ===== #
184         # YOUR CODE HERE:
185         # Initialize all parameters of the network in the self.params
dictionary.
186         # The weights and biases of layer 1 are W1 and b1; and in general the
187         # weights and biases of layer i are Wi and bi. The
188         # biases are initialized to zero and the weights are initialized
189         # so that each parameter has mean 0 and standard deviation
weight_scale.
190         # ===== #
191
192         dims = []
193         dims.append(input_dim)
194         dims.extend(hidden_dims)
195         dims.append(num_classes)
196
197         for i in np.arange(self.num_layers):
198             num = str(i+1)
199             self.params['W'+num] = weight_scale * np.random.randn(dims[i],
dims[i+1]) + 0
200             self.params['b'+num] = np.zeros((dims[i+1], 1))
201
202         # ===== #
203         # END YOUR CODE HERE
204         # ===== #
205
206         # When using dropout we need to pass a dropout_param dictionary to each
207         # dropout layer so that the layer knows the dropout probability and the
mode
208         # (train / test). You can pass the same dropout_param to each dropout
layer.
209         self.dropout_param = {}
210         if self.use_dropout:
211             self.dropout_param = {'mode': 'train', 'p': dropout}
212             if seed is not None:
213                 self.dropout_param['seed'] = seed
214
215         # With batch normalization we need to keep track of running means and
216         # variances, so we need to pass a special bn_param object to each batch
217         # normalization layer. You should pass self.bn_params[0] to the forward
pass
218         # of the first batch normalization layer, self.bn_params[1] to the
forward
219         # pass of the second batch normalization layer, etc.
220         self.bn_params = []
221         if self.use_batchnorm:

```

```

222     self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers
223 - 1)]
224
225     # Cast all parameters to the correct datatype
226     for k, v in self.params.items():
227         self.params[k] = v.astype(dtype)
228
229 def loss(self, X, y=None):
230     """
231     Compute loss and gradient for the fully-connected net.
232
233     Input / output: Same as TwoLayerNet above.
234     """
235     X = X.astype(self.dtype)
236     mode = 'test' if y is None else 'train'
237
238     # Set train/test mode for batchnorm params and dropout param since they
239     # behave differently during training and testing.
240     if self.dropout_param is not None:
241         self.dropout_param['mode'] = mode
242     if self.use_batchnorm:
243         for bn_param in self.bn_params:
244             bn_param[mode] = mode
245
246     scores = None
247
248     # ===== #
249     # YOUR CODE HERE:
250     #   Implement the forward pass of the FC net and store the output
251     #   scores as the variable "scores".
252     # ===== #
253
254     outs = {}
255     h = {}
256     h[0] = [X]
257
258     for i in np.arange(self.num_layers):
259         num = str(i+1)
260         outs[i+1] = affine_forward(h[i][0], self.params['W'+num],
self.params['b'+num])
261         if i != (self.num_layers-1):
262             h[i+1] = relu_forward(outs[i+1][0])
263
264     scores = outs[self.num_layers][0]
265
266     # ===== #
267     # END YOUR CODE HERE
268     # ===== #
269
270     # If test mode return early
271     if mode == 'test':
272         return scores
273
274     loss, grads = 0.0, {}
275     # ===== #
276     # YOUR CODE HERE:
277     #   Implement the backwards pass of the FC net and store the gradients
278     #   in the grads dict, so that grads[k] is the gradient of self.params[k]
279     #   Be sure your L2 regularization includes a 0.5 factor.

```

```
280 # ===== #
281
282 loss, dx = softmax_loss(scores, y)
283 reg_loss_sum = 0
284 for i in np.arange(self.num_layers):
285     num = str(i+1)
286     reg_loss_sum += np.linalg.norm(self.params['W'+num], 'fro')**2
287
288 loss += 0.5 * self.reg * reg_loss_sum
289
290 dict_dW = {}
291 dict_db = {}
292
293 dict_da = {}
294 dict_da[self.num_layers] = dx
295
296 for i in np.arange(self.num_layers, 0, -1):
297     dh, dW, db = affine_backward(dict_da[i], outs[i][1])
298     dict_dW[i] = dW
299     dict_db[i] = db
300
301     if i != 1:
302         dict_da[i-1] = relu_backward(dh, h[i-1][1])
303
304 for i in np.arange(self.num_layers):
305     num = str(i+1)
306     grads['W'+num] = dict_dW[i+1] + self.reg * self.params['W'+num]
307     grads['b'+num] = dict_db[i+1].T
308
309 # ===== #
310 # END YOUR CODE HERE
311 # ===== #
312 return loss, grads
313
```