

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4
5 class TwoLayerNet(object):
6     """
7     A two-layer fully-connected neural network. The net has an input dimension
8     of
9     N, a hidden layer dimension of H, and performs classification over C
10    classes.
11    We train the network with a softmax loss function and L2 regularization on
12    the
13    weight matrices. The network uses a ReLU nonlinearity after the first fully
14    connected layer.
15
16    In other words, the network has the following architecture:
17
18    input - fully connected layer - ReLU - fully connected layer - softmax
19
20    The outputs of the second fully-connected layer are the scores for each
21    class.
22    """
23
24    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
25        """
26        Initialize the model. Weights are initialized to small random values and
27        biases are initialized to zero. Weights and biases are stored in the
28        variable self.params, which is a dictionary with the following keys:
29
30        W1: First layer weights; has shape (H, D)
31        b1: First layer biases; has shape (H,)
32        W2: Second layer weights; has shape (C, H)
33        b2: Second layer biases; has shape (C,)
34
35        Inputs:
36        - input_size: The dimension D of the input data.
37        - hidden_size: The number of neurons H in the hidden layer.
38        - output_size: The number of classes C.
39        """
40        self.params = {}
41        self.params['W1'] = std * np.random.randn(hidden_size, input_size)
42        self.params['b1'] = np.zeros(hidden_size)
43        self.params['W2'] = std * np.random.randn(output_size, hidden_size)
44        self.params['b2'] = np.zeros(output_size)
45
46    def loss(self, X, y=None, reg=0.0):
47        """
48        Compute the loss and gradients for a two layer fully connected neural
49        network.
50
51        Inputs:
52        - X: Input data of shape (N, D). Each X[i] is a training sample.
53        - y: Vector of training labels. y[i] is the label for X[i], and each y[i]
54        is
55        an integer in the range 0 <= y[i] < C. This parameter is optional; if
56        it
57        is not passed then we only return scores, and if it is passed then we
58        instead return the loss and gradients.

```

```

54     - reg: Regularization strength.
55
56     Returns:
57     If y is None, return a matrix scores of shape (N, C) where scores[i, c]
is
58     the score for class c on input X[i].
59
60     If y is not None, instead return a tuple of:
61     - loss: Loss (data loss and regularization loss) for this batch of
training
62     samples.
63     - grads: Dictionary mapping parameter names to gradients of those
parameters
64     with respect to the loss function; has the same keys as self.params.
65     """
66     # Unpack variables from the params dictionary
67     W1, b1 = self.params['W1'], self.params['b1']
68     W2, b2 = self.params['W2'], self.params['b2']
69     N, D = X.shape
70
71     # Compute the forward pass
72     scores = None
73
74     # ===== #
75     # YOUR CODE HERE:
76     # Calculate the output scores of the neural network. The result
77     # should be (N, C). As stated in the description for this class,
78     # there should not be a ReLU layer after the second FC layer.
79     # The output of the second FC layer is the output scores. Do not
80     # use a for loop in your implementation.
81     # ===== #
82
83     relu = lambda x: np.maximum(x, 0)
84
85     h1 = np.dot(X, W1.T) + b1
86     scores = np.dot(relu(h1), W2.T) + b2
87
88     # ===== #
89     # END YOUR CODE HERE
90     # ===== #
91
92
93     # If the targets are not given then jump out, we're done
94     if y is None:
95         return scores
96
97     # Compute the loss
98     loss = None
99
100    # ===== #
101    # YOUR CODE HERE:
102    # Calculate the loss of the neural network. This includes the
103    # softmax loss and the L2 regularization for W1 and W2. Store the
104    # total loss in the variable loss. Multiply the regularization
105    # loss by 0.5 (in addition to the factor reg).
106    # ===== #
107
108    # scores is num_examples by num_classes
109
110    scores -= np.max(scores, axis=1, keepdims=True)

```

```

111     scores_exp = np.exp(scores)
112
113     probs = scores_exp / np.sum(scores_exp, axis=1, keepdims=True)
114     probs_row = probs[range(N), y]
115     probs_log = -np.log(probs_row)
116     softmax_loss = np.sum(probs_log) / N
117
118     reg_loss = 0.5 * reg * (np.linalg.norm(W1, 'fro')**2 + np.linalg.norm(W2,
'fro')**2)
119
120     loss = softmax_loss + reg_loss
121
122     # ===== #
123     # END YOUR CODE HERE
124     # ===== #
125
126     grads = {}
127
128     # ===== #
129     # YOUR CODE HERE:
130     # Implement the backward pass. Compute the derivatives of the
131     # weights and the biases. Store the results in the grads
132     # dictionary. e.g., grads['W1'] should store the gradient for
133     # W1, and be of the same size as W1.
134     # ===== #
135
136     probs[range(N), y] -= 1
137
138     dLdb = probs / N
139     dLdW2 = np.maximum(np.dot(W1, X.T) + b1.reshape([W1.shape[0], 1]), 0)
140
141     grads['W2'] = np.dot(dLdb.T, dLdW2.T) + reg * W2
142     grads['b2'] = np.sum(dLdb, axis=0, keepdims=True)
143
144     dbdh = W2.T
145     dLda = np.dot(dbdh, dLdb.T) * (np.dot(W1, X.T) > 0)
146
147     grads['W1'] = np.dot(dLda, X) + reg * W1
148     grads['b1'] = np.sum(dLda, axis=1, keepdims=True).T
149
150     # ===== #
151     # END YOUR CODE HERE
152     # ===== #
153
154     return loss, grads
155
156 def train(self, X, y, X_val, y_val,
157           learning_rate=1e-3, learning_rate_decay=0.95,
158           reg=1e-5, num_iters=100,
159           batch_size=200, verbose=False):
160     """
161     Train this neural network using stochastic gradient descent.
162
163     Inputs:
164     - X: A numpy array of shape (N, D) giving training data.
165     - y: A numpy array of shape (N,) giving training labels; y[i] = c means
that X[i] has label c, where 0 <= c < C.
166     - X_val: A numpy array of shape (N_val, D) giving validation data.
167     - y_val: A numpy array of shape (N_val,) giving validation labels.

```

```

169     - learning_rate: Scalar giving learning rate for optimization.
170     - learning_rate_decay: Scalar giving factor used to decay the learning
rate
171         after each epoch.
172     - reg: Scalar giving regularization strength.
173     - num_iters: Number of steps to take when optimizing.
174     - batch_size: Number of training examples to use per step.
175     - verbose: boolean; if true print progress during optimization.
176     """
177     num_train = X.shape[0]
178     iterations_per_epoch = max(num_train / batch_size, 1)
179
180     # Use SGD to optimize the parameters in self.model
181     loss_history = []
182     train_acc_history = []
183     val_acc_history = []
184
185     for it in np.arange(num_iters):
186         X_batch = None
187         y_batch = None
188
189         # ===== #
190         # YOUR CODE HERE:
191         #   Create a minibatch by sampling batch_size samples randomly.
192         # ===== #
193
194         idx = np.random.choice(num_train, batch_size)
195         X_batch = X[idx]
196         y_batch = y[idx]
197
198         # ===== #
199         # END YOUR CODE HERE
200         # ===== #
201
202         # Compute loss and gradients using the current minibatch
203         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
204         loss_history.append(loss)
205
206         # ===== #
207         # YOUR CODE HERE:
208         #   Perform a gradient descent step using the minibatch to update
209         #   all parameters (i.e., W1, W2, b1, and b2).
210         # ===== #
211
212         self.params['W1'] = self.params['W1'] - learning_rate * grads['W1']
213         self.params['b1'] = self.params['b1'] - learning_rate * grads['b1']
214
215         self.params['W2'] = self.params['W2'] - learning_rate * grads['W2']
216         self.params['b2'] = self.params['b2'] - learning_rate * grads['b2']
217
218         # ===== #
219         # END YOUR CODE HERE
220         # ===== #
221
222         if verbose and it % 100 == 0:
223             print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
224
225         # Every epoch, check train and val accuracy and decay learning rate.
226         if it % iterations_per_epoch == 0:
227             # Check accuracy

```

```

228     train_acc = (self.predict(X_batch) == y_batch).mean()
229     val_acc = (self.predict(X_val) == y_val).mean()
230     train_acc_history.append(train_acc)
231     val_acc_history.append(val_acc)
232
233     # Decay learning rate
234     learning_rate *= learning_rate_decay
235
236     return {
237         'loss_history': loss_history,
238         'train_acc_history': train_acc_history,
239         'val_acc_history': val_acc_history,
240     }
241
242 def predict(self, X):
243     """
244     Use the trained weights of this two-layer network to predict labels for
245     data points. For each data point we predict scores for each of the C
246     classes, and assign each data point to the class with the highest score.
247
248     Inputs:
249     - X: A numpy array of shape (N, D) giving N D-dimensional data points to
250         classify.
251
252     Returns:
253     - y_pred: A numpy array of shape (N,) giving predicted labels for each of
254         the elements of X. For all i, y_pred[i] = c means that X[i] is
255         predicted
256         to have class c, where 0 <= c < C.
257     """
258     y_pred = None
259
260     # ===== #
261     # YOUR CODE HERE:
262     # Predict the class given the input data.
263     # ===== #
264
265     relu = lambda x: np.maximum(x, 0)
266
267     h1 = np.dot(X, self.params['W1'].T) + self.params['b1']
268     scores = np.dot(relu(h1), self.params['W2'].T) + self.params['b2']
269
270     y_pred = np.argmax(scores, axis=1)
271
272     # ===== #
273     # END YOUR CODE HERE
274     # ===== #
275
276     return y_pred

```