

```

1 import numpy as np
2 from .layers import *
3 from .layer_utils import *
4
5
6 class TwoLayerNet(object):
7     """
8     A two-layer fully-connected neural network with ReLU nonlinearity and
9     softmax loss that uses a modular layer design. We assume an input dimension
10    of D, a hidden dimension of H, and perform classification over C classes.
11
12    The architecture should be affine - relu - affine - softmax.
13
14    Note that this class does not implement gradient descent; instead, it
15    will interact with a separate Solver object that is responsible for running
16    optimization.
17
18    The learnable parameters of the model are stored in the dictionary
19    self.params that maps parameter names to numpy arrays.
20    """
21
22    def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
23                  dropout=0, weight_scale=1e-3, reg=0.0):
24        """
25        Initialize a new network.
26
27        Inputs:
28        - input_dim: An integer giving the size of the input
29        - hidden_dims: An integer giving the size of the hidden layer
30        - num_classes: An integer giving the number of classes to classify
31        - dropout: Scalar between 0 and 1 giving dropout strength.
32        - weight_scale: Scalar giving the standard deviation for random
33          initialization of the weights.
34        - reg: Scalar giving L2 regularization strength.
35        """
36        self.params = {}
37        self.reg = reg
38
39        # ===== #
40        # YOUR CODE HERE:
41        # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
42        # self.params['W2'], self.params['b1'] and self.params['b2']. The
43        # biases are initialized to zero and the weights are initialized
44        # so that each parameter has mean 0 and standard deviation
45        weight_scale.
46        # The dimensions of W1 should be (input_dim, hidden_dim) and the
47        # dimensions of W2 should be (hidden_dims, num_classes)
48        # ===== #
49        self.params['W1'] = weight_scale * np.random.randn(input_dim,
50        hidden_dims) + 0
51        self.params['W2'] = weight_scale * np.random.randn(hidden_dims,
52        num_classes) + 0
53
54        self.params['b1'] = np.zeros((hidden_dims, 1))
55        self.params['b2'] = np.zeros((num_classes, 1))
56
57        # ===== #
58        # END YOUR CODE HERE

```

```

57     # ===== #
58
59     def loss(self, X, y=None):
60         """
61         Compute loss and gradient for a minibatch of data.
62
63         Inputs:
64         - X: Array of input data of shape (N, d_1, ..., d_k)
65         - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
66
67         Returns:
68         If y is None, then run a test-time forward pass of the model and return:
69         - scores: Array of shape (N, C) giving classification scores, where
70           scores[i, c] is the classification score for X[i] and class c.
71
72         If y is not None, then run a training-time forward and backward pass and
73         return a tuple of:
74         - loss: Scalar value giving the loss
75         - grads: Dictionary with the same keys as self.params, mapping parameter
76           names to gradients of the loss with respect to those parameters.
77         """
78         scores = None
79
80         # ===== #
81         # YOUR CODE HERE:
82         #   Implement the forward pass of the two-layer neural network. Store
83         #   the class scores as the variable 'scores'. Be sure to use the layers
84         #   you prior implemented.
85         # ===== #
86
87         out_l1, cache_l1 = affine_forward(X, self.params['W1'],
88 self.params['b1'])
89         out_relu, cache_relu = relu_forward(out_l1)
90         scores, cache_l2 = affine_forward(out_relu, self.params['W2'],
91 self.params['b2'])
92
93         # ===== #
94         # END YOUR CODE HERE
95         # ===== #
96
97         # If y is None then we are in test mode so just return scores
98         if y is None:
99             return scores
100
101         loss, grads = 0, {}
102
103         # ===== #
104         # YOUR CODE HERE:
105         #   Implement the backward pass of the two-layer neural net. Store
106         #   the loss as the variable 'loss' and store the gradients in the
107         #   'grads' dictionary. For the grads dictionary, grads['W1'] holds
108         #   the gradient for W1, grads['b1'] holds the gradient for b1, etc.
109         #   i.e., grads[k] holds the gradient for self.params[k].
110         #
111         #   Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
112         #   for each W. Be sure to include the 0.5 multiplying factor to
113         #   match our implementation.
114         #
115         #   And be sure to use the layers you prior implemented.
116         # ===== #

```

```

115     loss, dx2 = softmax_loss(scores, y)
116     reg_loss = 0.5 * self.reg * (np.linalg.norm(self.params['W1'], 'fro')**2
+ np.linalg.norm(self.params['W2'], 'fro')**2)
117     loss += reg_loss
118
119     dh1, dw2, db2 = affine_backward(dx2, cache_l2)
120     da = relu_backward(dh1, cache_relu)
121     dx1, dw1, db1 = affine_backward(da, cache_l1)
122
123     grads['W1'] = dw1 + self.reg * self.params['W1']
124     grads['b1'] = db1.T
125
126     grads['W2'] = dw2 + self.reg * self.params['W2']
127     grads['b2'] = db2.T
128
129     # ===== #
130     # END YOUR CODE HERE
131     # ===== #
132
133     return loss, grads
134
135
136 class FullyConnectedNet(object):
137     """
138     A fully-connected neural network with an arbitrary number of hidden layers,
139     ReLU nonlinearities, and a softmax loss function. This will also implement
140     dropout and batch normalization as options. For a network with L layers,
141     the architecture will be
142
143     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
144
145     where batch normalization and dropout are optional, and the {...} block is
146     repeated L - 1 times.
147
148     Similar to the TwoLayerNet above, learnable parameters are stored in the
149     self.params dictionary and will be learned using the Solver class.
150     """
151
152     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
153                 dropout=0, use_batchnorm=False, reg=0.0,
154                 weight_scale=1e-2, dtype=np.float32, seed=None):
155         """
156         Initialize a new FullyConnectedNet.
157
158         Inputs:
159         - hidden_dims: A list of integers giving the size of each hidden layer.
160         - input_dim: An integer giving the size of the input.
161         - num_classes: An integer giving the number of classes to classify.
162         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0
163         then
164             the network should not use dropout at all.
165         - use_batchnorm: Whether or not the network should use batch
166         normalization.
167         - reg: Scalar giving L2 regularization strength.
168         - weight_scale: Scalar giving the standard deviation for random
169         initialization of the weights.
170         - dtype: A numpy datatype object; all computations will be performed
171         using
172             this datatype. float32 is faster but less accurate, so you should use
173             float64 for numeric gradient checking.

```

```

171     - seed: If not None, then pass this random seed to the dropout layers.
This
172     will make the dropout layers deterministic so we can gradient check the
173     model.
174     """
175     self.use_batchnorm = use_batchnorm
176     self.use_dropout = dropout > 0
177     self.reg = reg
178     self.num_layers = 1 + len(hidden_dims)
179     self.dtype = dtype
180     self.params = {}
181
182     # ===== #
183     # YOUR CODE HERE:
184     # Initialize all parameters of the network in the self.params
dictionary.
185     # The weights and biases of layer 1 are W1 and b1; and in general the
186     # weights and biases of layer i are Wi and bi. The
187     # biases are initialized to zero and the weights are initialized
188     # so that each parameter has mean 0 and standard deviation
weight_scale.
189     #
190     # BATCHNORM: Initialize the gammas of each layer to 1 and the beta
191     # parameters to zero. The gamma and beta parameters for layer 1 should
192     # be self.params['gamma1'] and self.params['beta1']. For layer 2, they
193     # should be gamma2 and beta2, etc. Only use batchnorm if
self.use_batchnorm
194     # is true and DO NOT do batch normalize the output scores.
195     # ===== #
196
197     dims = []
198     dims.append(input_dim)
199     dims.extend(hidden_dims)
200     dims.append(num_classes)
201
202     for i in np.arange(self.num_layers):
203         num = str(i+1)
204         self.params['W'+num] = weight_scale * np.random.randn(dims[i],
dims[i+1]) + 0
205         self.params['b'+num] = np.zeros((dims[i+1]))
206
207         if i == (self.num_layers - 1):
208             break
209
210         if self.use_batchnorm:
211             self.params['gamma'+num] = np.ones((dims[i+1]))
212             self.params['beta'+num] = np.zeros((dims[i+1]))
213
214     # ===== #
215     # END YOUR CODE HERE
216     # ===== #
217
218     # When using dropout we need to pass a dropout_param dictionary to each
219     # dropout layer so that the layer knows the dropout probability and the
mode
220     # (train / test). You can pass the same dropout_param to each dropout
layer.
221     self.dropout_param = {}
222     if self.use_dropout:
223         self.dropout_param = {'mode': 'train', 'p': dropout}

```

```

224         if seed is not None:
225             self.dropout_param['seed'] = seed
226
227         # With batch normalization we need to keep track of running means and
228         # variances, so we need to pass a special bn_param object to each batch
229         # normalization layer. You should pass self.bn_params[0] to the forward
pass
230         # of the first batch normalization layer, self.bn_params[1] to the
forward
231         # pass of the second batch normalization layer, etc.
232         self.bn_params = []
233         if self.use_batchnorm:
234             self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers
- 1)]
235
236         # Cast all parameters to the correct datatype
237         for k, v in self.params.items():
238             self.params[k] = v.astype(dtype)
239
240
241     def loss(self, X, y=None):
242         """
243         Compute loss and gradient for the fully-connected net.
244
245         Input / output: Same as TwoLayerNet above.
246         """
247         X = X.astype(self.dtype)
248         mode = 'test' if y is None else 'train'
249
250         # Set train/test mode for batchnorm params and dropout param since they
251         # behave differently during training and testing.
252         if self.dropout_param is not None:
253             self.dropout_param['mode'] = mode
254         if self.use_batchnorm:
255             for bn_param in self.bn_params:
256                 bn_param['mode'] = mode
257
258         scores = None
259
260         # ===== #
261         # YOUR CODE HERE:
262         # Implement the forward pass of the FC net and store the output
263         # scores as the variable "scores".
264         #
265         # BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
266         # between the affine_forward and relu_forward layers. You may
267         # also write an affine_batchnorm_relu() function in layer_utils.py.
268         #
269         # DROPOUT: If dropout is non-zero, insert a dropout layer after
270         # every ReLU layer.
271         # ===== #
272
273         fc_outs = {}
274         fc_caches = {}
275
276         batchnorm_outs = {}
277         batchnorm_caches = {}
278
279         h = {}
280         h[0] = X

```

```

281     relu_caches = {}
282
283     dropout_caches = {}
284
285     for i in np.arange(self.num_layers):
286         num = str(i+1)
287         # fc
288         fc_outs[i+1], fc_caches[i+1] = affine_forward(h[i],
self.params['W'+num], self.params['b'+num])
289         if i == (self.num_layers - 1):
290             break
291         relu_input = fc_outs[i+1]
292
293         # batch-norm
294         if self.use_batchnorm:
295             batchnorm_outs[i+1], batchnorm_caches[i+1] =
batchnorm_forward(fc_outs[i+1], self.params['gamma'+num],
self.params['beta'+num], self.bn_params[i])
296             relu_input = batchnorm_outs[i+1]
297
298         # relu
299         h[i+1], relu_caches[i+1] = relu_forward(relu_input)
300
301         # dropout
302         if self.use_dropout:
303             h[i+1], dropout_caches[i+1] = dropout_forward(h[i+1],
self.dropout_param)
304
305     scores = fc_outs[self.num_layers]
306
307     # ===== #
308     # END YOUR CODE HERE
309     # ===== #
310
311     # If test mode return early
312     if mode == 'test':
313         return scores
314
315     loss, grads = 0.0, {}
316     # ===== #
317     # YOUR CODE HERE:
318     # Implement the backwards pass of the FC net and store the gradients
319     # in the grads dict, so that grads[k] is the gradient of self.params[k]
320     # Be sure your L2 regularization includes a 0.5 factor.
321     #
322     # BATCHNORM: Incorporate the backward pass of the batchnorm.
323     #
324     # DROPOUT: Incorporate the backward pass of dropout.
325     # ===== #
326
327     loss, dx = softmax_loss(scores, y)
328     reg_loss_sum = 0
329     for i in np.arange(self.num_layers):
330         num = str(i+1)
331         reg_loss_sum += np.linalg.norm(self.params['W'+num], 'fro')**2
332     loss += 0.5 * self.reg * reg_loss_sum
333
334     dict_dW = {}
335     dict_db = {}
336

```

```

337     dict_dgamma = {}
338     dict_dbeta = {}
339
340     # fc - last layer
341     dh, dW, db = affine_backward(dx, fc_caches[self.num_layers])
342     dict_dW[self.num_layers] = dW
343     dict_db[self.num_layers] = db
344
345     for i in np.arange(self.num_layers - 1, 0, -1):
346         # dropout
347         if self.use_dropout:
348             dh = dropout_backward(dh, dropout_caches[i])
349
350         # relu
351         dx_relu = relu_backward(dh, relu_caches[i])
352         fc_input = dx_relu
353
354         # batch-norm
355         if self.use_batchnorm:
356             dx_batchnorm, dgamma, dbeta = batchnorm_backward(dx_relu,
batchnorm_caches[i])
357             dict_dgamma[i] = dgamma
358             dict_dbeta[i] = dbeta
359             fc_input = dx_batchnorm
360
361         # fc
362         dh, dW, db = affine_backward(fc_input, fc_caches[i])
363         dict_dW[i] = dW
364         dict_db[i] = db
365
366     for i in np.arange(self.num_layers):
367         num = str(i+1)
368         grads['W'+num] = dict_dW[i+1] + self.reg * self.params['W'+num]
369         grads['b'+num] = dict_db[i+1]
370         if i == (self.num_layers - 1):
371             break
372         if self.use_batchnorm:
373             grads['gamma'+num] = dict_dgamma[i+1]
374             grads['beta'+num] = dict_dbeta[i+1]
375
376     # ===== #
377     # END YOUR CODE HERE
378     # ===== #
379
380     return loss, grads
381

```